

Jeremy McMahan

✉ jeremymcmahan@gmail.com • 🌐 www.jeremymcmahan.com • 🌀 [jermcmahan](#)
[in jermcmahan](#) • 📺 [m8robxcAAAAJ](#)

Research Overview

Developing algorithmic foundations for decision-making under uncertainty. My work bridges *machine learning* and *worst-case algorithm design* to create provably efficient, robust, and constraint-satisfying agents able to solve complex combinatorial and multi-agent tasks — overcoming intractability barriers to ensure capable and reliable AI agents.

Education

- 2018–2025 **PhD in Computer Science**, *University of Wisconsin-Madison*
- Thesis: Safe Multi-Agent Reinforcement Learning in Polynomial Time
 - Advisor: Prof. Jerry Zhu
 - **LEAP Fellow, IFDS Research Position**
- 2018–2020 **MS in Computer Science**, *University of Wisconsin-Madison*
- Thesis: A D-competitive algorithm for the Multilevel Aggregation Problem with Deadlines
 - Advisor: Prof. Shuchi Chawla
- 2016–2018 **BS in Mathematics and Computer Science**, *University of Illinois Urbana-Champaign*
- Thesis: Spectral Graph Isomorphism
 - Advisor: Alexandra Kolla
 - **Student ranked #1 in Mathematics and Computer Science**

Core Research Contributions

Safe and Constrained Agent Planning: Designed the *first polynomial-time approximation algorithms for general constrained MDPs*, overcoming long-standing intractability barriers in constrained decision-making. Enables constraint-aware agents that satisfy hard operational and safety constraints under uncertainty. (*ICML '25, NeurIPS '24, AISTATS '24*)

Robust Multi-Agent Decision-Making: Developed the *first tractable algorithms for computing equilibria robust to model uncertainty and environmental manipulation* in Markov games. Enables incentive-aware agents that operate robustly in adversarial and strategically-complex multi-agent environments. (*ICML '25, ICML '24, AAAI '24*)

Behavioral Influence in Multi-Agent Systems: *Characterized the capabilities and limits of behavioral influence in Markov games* — what behaviors an influencer can tractably induce through strategic manipulation or instruction. Provides tools for auditing, securing, and aligning the behavior of rational agents. (*ICML '24, RLC '24, AAAI '24, AAAI '23*)

Online and Adaptive Optimization: Established the state-of-the-art competitive and adaptive approximation guarantees for multi-level aggregation and correlated Pandora's Box, advancing tractable optimization for large-scale decision-making under uncertainty. (*APPROX '23*)

Publications

**Indicates equal contribution. Authorship order follows venue conventions.*

- ICML '25 **Polynomial-Time Approximability of Constrained Reinforcement Learning**
Jeremy McMahan
- ICML '25 **Anytime-Constrained Equilibria in Polynomial Time**
Jeremy McMahan
- NeurIPS '24 **Deterministic Policies for Constrained Reinforcement Learning in Polynomial Time**
Jeremy McMahan
- RLC '24 **Inception: Efficiently Computable Misinformation Attacks on Markov Games**
Jeremy McMahan, Young Wu, Yudong Chen, Xiaojin Zhu, Qiaomin Xie [Oral]
- ICML '24 **Roping in Uncertainty: Robustness and Regularization in Markov Games**
Jeremy McMahan, Giovanni Artiglio, Qiaomin Xie
- ICML '24 **Minimally Modifying a Markov Game to Achieve any Nash Equilibrium and Value**
Young Wu, Jeremy McMahan, Yiding Chen, Yudong Chen, Xiaojin Zhu, Qiaomin Xie
- CogSci '24 **Various Misleading Visual Features in Misleading Graphs: Do they truly deceive us?**
Jihyun Rho, Martina A Rau, Shubham Kumar Bharti, Rosanne Luu, Jeremy McMahan, Andrew Wang, Jerry Zhu
- AISTATS '24 **Anytime-Constrained Reinforcement Learning**
Jeremy McMahan, Xiaojin Zhu
- AAAI '24 **Optimal Attack and Defense for Reinforcement Learning**
Jeremy McMahan, Young Wu, Xiaojin Zhu, Qiaomin Xie [Oral]
- AAAI '24 **Data Poisoning to Fake a Nash Equilibrium for Markov Games**
Young Wu, Jeremy McMahan, Xiaojin Zhu, Qiaomin Xie
- APPROX '23 **Approximating Pandora's Box with Correlations**
Shuchi Chawla, Evangelia Gergatsouli*, Jeremy McMahan*, Christos Tzamos* [Oral]*
- AAAI '23 **Reward Poisoning Attacks on Offline Multi-Agent Reinforcement Learning**
Young Wu, Jeremy McMahan, Xiaojin Zhu, Qiaomin Xie [Oral]
- GAMES '20 **Noble Deceit: Optimizing Social Welfare for Myopic Multi-Armed Bandits**
Ashwin Maran, Jeremy McMahan*, Nathaniel Sauerberg* [Oral]*

Research Software

LLM-Augmented Optimization Agent, [\[GitHub Link\]](#)

An LLM-augmented agent architecture for combinatorial optimization that integrates LLM-generated heuristics with classical algorithm design — achieving empirical efficiency gains from experience while provably retaining worst-case optimality guarantees, demonstrating a principled path to integrating LLM reasoning with formal algorithmic guarantees.

Constrained Sequential Decision-Making, [\[GitHub Link\]](#)

An algorithmic framework for constrained decision-making, implementing efficient optimization solvers and deep reinforcement learning approaches for anytime-constrained agents. Includes a comprehensive benchmarking suite for evaluating performance tradeoffs. (*AISTATS '24*)

Robust Multi-Agent Decision-Making, [\[GitHub Link\]](#)

Reference implementations of scalable algorithms for online attacks and defenses in Markov games, enabling empirical evaluation and equilibrium computation for robust multi-agent decision systems. (*AAAI '24*)

Optimization and Learning for Resource Allocation, [\[GitHub Link\]](#)

A large-scale evaluation framework demonstrating the tradeoffs between classical approximation algorithms and deep-learning heuristics on resource allocation tasks, providing empirical grounding for theory-first algorithm design in AI systems.

Selected Talks

From Knapsacks to Self Driving: FPTAS Recipes for Constrained Reinforcement Learning

- UMD College Park (MARL Seminar) *March 2025*
- UW-Madison (Theory of Computing Seminar) *November 2024*

Optimal Attack and Defense for Reinforcement Learning

- UW-Madison (Theory of Computing Seminar) *February 2025*
- UW-Madison (Computer Vision Roundtable) *January 2024*

Teaching Experience

Primary Instructor

Summer 2022 **CS 577: Introduction to Algorithms**, *University of Wisconsin-Madison*

- Responsible for curriculum design, lecture delivery, course structure, and student evaluation for a class of 100 students.
- Designed original lecture slides, lecture notes, lecture videos, and problem sets focusing on the design of efficient algorithms and rigorous proofs of correctness.
- Managed one TA and grader.

Head Teaching Assistant

Spring 2021, 2022 **CS 540: Introduction to Artificial Intelligence**, *University of Wisconsin-Madison*

- Coordinated the administrative and grading workflow for a large-scale course with 500 students, 8 TAs, and over 20 graders and peer mentors.
- Developed standardized grading rubrics and course structures in addition to the standard teaching assistant duties.
- Instructors: Ilias Diakonikolas, Sharon Li, Fred Sala, and Jerry Zhu

Guest Lecturer

- Spring 2025 **CS 839: Game Theory, Optimization & Learning**, *University of Wisconsin-Madison*
- Designed and presented two lectures on the theory of multi-agent reinforcement learning.
 - Instructor: Manolis Vlatakis
- Summer 2021, 2022 **New Horizons in Theoretical Computer Science Summer School**, *ACM SIGACT*
- Designed lectures spanning complexity theory basics, interactive proofs, and Markov-chain Monte Carlo methods.
 - Instructors: Antonio Blanca, Nicole Immorlica, Yael Kalai, Jelani Nelson, and Ashia Wilson
- Graduate Teaching Assistant**
- Fall 2020, 2021 **CS 540: Introduction to Artificial Intelligence**, *University of Wisconsin-Madison*
Instructors: Anthony Gitter, Josiah Hanna, Yin Li, Yingyu Liang, and Daifeng Wang
- Spring 2019, 2020 **CS 577: Introduction to Algorithms**, *University of Wisconsin-Madison*
Instructors: Shuchi Chawla, Dieter van Melkebeek, and Christos Tzamos
- Fall 2019 **CS 301: Introduction to Data Science**, *University of Wisconsin-Madison*
Instructor: Tyler Caraza-Harter
- Fall 2018 **CS 524: Introduction to Optimization**, *University of Wisconsin-Madison*
Instructor: Steve Wright

Awards and Honors

- 2021–2025 **LEAP Fellow**, *CMD-IT / NSF Alliance*
Selected for the NSF-funded Leadership, Excellence, Access, and Pathways (LEAP) Alliance fellowship, focused on diversifying future faculty leadership in computing.
- 2023 **Graduate Instructor Award**, *University of Wisconsin-Madison*
Recognition for excellent performance as a primary instructor.
- 2019–2022 **The "Golden Brick" Teaching Award**, *University of Wisconsin-Madison*
A student/instructor-voted honor for most effective teaching. Won consecutively for every year served as a TA (4 years).
- 2018–2019 **CS Departmental Scholarship**, *University of Wisconsin-Madison*
Merit-based scholarship awarded to incoming PhD students.
- 2018 **Most Outstanding Major Award**, *University of Illinois Urbana-Champaign*
Awarded to the #1-ranked student in Mathematics and Computer Science.

Professional Service

- 2024–2025 **Conference Reviewer**
NeurIPS (2024, 2025), ICML (2024, 2025), AISTATS (2025), RLC (2024)

Leadership and Organization

- 2023–2025 **Co-Organizer**, *Machine Learning Lunch Meetings (MLLM) Seminar*, *University of Wisconsin-Madison*
Organized weekly research seminars for the ML community, which included managing speaker invitations, scheduling, and logistics for internal and external presenters.

Mentorship and Outreach

- 2021–2022 **Planning Committee & Teaching Assistant**, *ACM SIGACT, New Horizons in Theoretical Computer Science Summer School*
A week-long summer school aimed at inspiring undergraduates from underrepresented backgrounds to pursue careers in TCS.
- Served on the planning committee.
 - Mentored students during technical sessions and social events.
 - Developed and presented lessons on selected TCS topics.

Technical Skills

- Areas: Agentic AI; Learning-Augmented Algorithms; Reinforcement Learning; Algorithms under Uncertainty; Algorithmic Game Theory; Approximation Algorithms; Online, Stochastic, and Combinatorial Optimization
- Languages: Python, Julia, C/C++, Mathematica, SageMath, $\LaTeX$, Bash
- Libraries: Gymnasium, LangGraph, PyTorch, OpenAI API, Gurobi, SciPy, NumPy, Pandas